

# Data 622 Assignment 1

Keith DeNivo

## Read in Data

```
file_url <- "https://raw.githubusercontent.com/division-zero/Data-622/refs/heads/main/assignment-1.csv"
# Download the file to a temporary location
temp_file <- tempfile(fileext = ".csv")
download.file(file_url, destfile = temp_file, mode = "wb")
# Read the csv file
fulldata <- read.delim(temp_file, sep = ";", header = TRUE, stringsAsFactors = FALSE)
# View the data
head(fulldata)
```

	age	job	marital	education	default	housing	loan	contact	month
1	56	housemaid	married	basic.4y	no	no	no	telephone	may
2	57	services	married	high.school	unknown	no	no	telephone	may
3	37	services	married	high.school	no	yes	no	telephone	may
4	40	admin.	married	basic.6y	no	no	no	telephone	may
5	56	services	married	high.school	no	no	yes	telephone	may
6	45	services	married	basic.9y	unknown	no	no	telephone	may
	day_of_week	duration	campaign	pdays	previous	poutcome	emp.var.rate		
1	mon	261	1	999	0	nonexistent		1.1	
2	mon	149	1	999	0	nonexistent		1.1	
3	mon	226	1	999	0	nonexistent		1.1	
4	mon	151	1	999	0	nonexistent		1.1	
5	mon	307	1	999	0	nonexistent		1.1	
6	mon	198	1	999	0	nonexistent		1.1	
	cons.price.idx	cons.conf.idx	euribor3m	nr.employed	y				
1	93.994	-36.4	4.857	5191	no				
2	93.994	-36.4	4.857	5191	no				

```

3      93.994      -36.4      4.857      5191 no
4      93.994      -36.4      4.857      5191 no
5      93.994      -36.4      4.857      5191 no
6      93.994      -36.4      4.857      5191 no

```

```

# Clean up the temporary file
unlink(temp_file)

```

```

#smaller data set for training or testing
file_url <- "https://raw.githubusercontent.com/division-zero/Data-622/refs/heads/main/assignm
# Download the file to a temporary location
temp_file <- tempfile(fileext = ".csv")
download.file(file_url, destfile = temp_file, mode = "wb")
# Read the csv file
partdata <- read.delim(file_url, sep = ";", header = TRUE, stringsAsFactors = FALSE)
# View the data
head(partdata)

```

	age	job	marital	education	default	housing	loan	contact
1	30	blue-collar	married	basic.9y	no	yes	no	cellular
2	39	services	single	high.school	no	no	no	telephone
3	25	services	married	high.school	no	yes	no	telephone
4	38	services	married	basic.9y	no	unknown	unknown	telephone
5	47	admin.	married	university.degree	no	yes	no	cellular
6	32	services	single	university.degree	no	no	no	cellular
	month	day_of_week	duration	campaign	pdays	previous	poutcome	emp.var.rate
1	may	fri	487	2	999	0	nonexistent	-1.8
2	may	fri	346	4	999	0	nonexistent	1.1
3	jun	wed	227	1	999	0	nonexistent	1.4
4	jun	fri	17	3	999	0	nonexistent	1.4
5	nov	mon	58	1	999	0	nonexistent	-0.1
6	sep	thu	128	3	999	2	failure	-1.1
	cons.price.idx	cons.conf.idx	euribor3m	nr.employed	y			
1	92.893	-46.2	1.313	5099.1	no			
2	93.994	-36.4	4.855	5191.0	no			
3	94.465	-41.8	4.962	5228.1	no			
4	94.465	-41.8	4.959	5228.1	no			
5	93.200	-42.0	4.191	5195.8	no			
6	94.199	-37.5	0.884	4963.6	no			

```
# Clean up the temporary file
unlink(temp_file)
```

## Basic data info

```
summary(fulldata)
```

age	job	marital	education
Min. :17.00	Length:41188	Length:41188	Length:41188
1st Qu.:32.00	Class :character	Class :character	Class :character
Median :38.00	Mode :character	Mode :character	Mode :character
Mean :40.02			
3rd Qu.:47.00			
Max. :98.00			
default	housing	loan	contact
Length:41188	Length:41188	Length:41188	Length:41188
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character
month	day_of_week	duration	campaign
Length:41188	Length:41188	Min. : 0.0	Min. : 1.000
Class :character	Class :character	1st Qu.: 102.0	1st Qu.: 1.000
Mode :character	Mode :character	Median : 180.0	Median : 2.000
		Mean : 258.3	Mean : 2.568
		3rd Qu.: 319.0	3rd Qu.: 3.000
		Max. :4918.0	Max. :56.000
pdays	previous	poutcome	emp.var.rate
Min. : 0.0	Min. :0.000	Length:41188	Min. :-3.40000
1st Qu.:999.0	1st Qu.:0.000	Class :character	1st Qu.:-1.80000
Median :999.0	Median :0.000	Mode :character	Median : 1.10000
Mean :962.5	Mean :0.173		Mean : 0.08189
3rd Qu.:999.0	3rd Qu.:0.000		3rd Qu.: 1.40000
Max. :999.0	Max. :7.000		Max. : 1.40000
cons.price.idx	cons.conf.idx	euribor3m	nr.employed
Min. :92.20	Min. :-50.8	Min. :0.634	Min. :4964
1st Qu.:93.08	1st Qu.:-42.7	1st Qu.:1.344	1st Qu.:5099
Median :93.75	Median :-41.8	Median :4.857	Median :5191
Mean :93.58	Mean :-40.5	Mean :3.621	Mean :5167

```

3rd Qu.:93.99   3rd Qu.:-36.4   3rd Qu.:4.961   3rd Qu.:5228
Max.    :94.77   Max.    :-26.9   Max.    :5.045   Max.    :5228
y
Length:41188
Class :character
Mode  :character

```

#converting yes to 1 and no to 0 for binary classification

```

partdata$term_deposit <- ifelse(partdata$y == "yes", 1, 0)
partdata$term_deposit_factor <- factor(partdata$term_deposit, levels = c(0, 1))

```

## Base models

### Build a tree model based on full data set

cp complexity parameter default is 0.01

```

fulldata$term_deposit <- ifelse(fulldata$y == "yes", 1, 0)
fulldata$term_deposit_factor <- factor(fulldata$term_deposit, levels = c(0, 1))
set.seed(42)
#removing columns/features that are directly related to target variable
modedata <- fulldata |> dplyr::select(-y, -term_deposit, -duration)
#taking the full data set and splitting 0.7:0.3 train:test split.
train_index <- sample(seq_len(nrow(modedata)), size = 0.7 * nrow(modedata))
train_data <- modedata[train_index, ]
test_data <- modedata[-train_index, ]

#train on full data set without missing features
tree_model <- rpart(
  term_deposit_factor ~ .,
  data = train_data,
  method = "class",
  control = rpart.control(
    cp = 0.01,
    minsplit = 20,
    maxdepth = 30


```

```
)  
)  
  
tree_model
```

```
n= 28831
```

```
node), split, n, loss, yval, (yprob)  
* denotes terminal node  
  
1) root 28831 3236 0 (0.88775970 0.11224030)  
2) nr.employed>=5087.65 25369 1724 0 (0.93204304 0.06795696) *  
3) nr.employed< 5087.65 3462 1512 0 (0.56325823 0.43674177)  
6) pdays>=513 2599 909 0 (0.65025010 0.34974990) *  
7) pdays< 513 863 260 1 (0.30127462 0.69872538) *
```

```
rpart.plot(tree_model)
```



```
pred_class <- predict(tree_model, test_data, type = "class")
```

```
table(Predicted = pred_class, Actual = test_data$term_deposit_factor)
```

		Actual
Predicted	0	1
0	10841	1134
1	112	270

```
mean(pred_class == test_data$term_deposit_factor)
```

[1] 0.8991665

```
confusionMatrix(pred_class, test_data$term_deposit_factor)
```

#### Confusion Matrix and Statistics

		Reference
Prediction	0	1
0	10841	1134
1	112	270

Accuracy : 0.8992  
95% CI : (0.8937, 0.9044)

No Information Rate : 0.8864  
P-Value [Acc > NIR] : 2.847e-06

Kappa : 0.2667

Mcnemar's Test P-Value : < 2.2e-16

Sensitivity : 0.9898  
Specificity : 0.1923  
Pos Pred Value : 0.9053  
Neg Pred Value : 0.7068  
Prevalence : 0.8864  
Detection Rate : 0.8773  
Detection Prevalence : 0.9691  
Balanced Accuracy : 0.5910

'Positive' Class : 0

```
Accuracy : 0.8992
```

```
sensitivity 0.9898
```

```
specificity 0.1923
```

```
kappa 0.2667
```

## Random Forest

```
#random forest model
set.seed(42)
rf_model <- randomForest(
  term_deposit_factor ~ .,
  data = train_data,
  ntree = 500,
  importance = TRUE
)

forrest_pred_class <- predict(rf_model, test_data, type = "class")

rf_model
```

Call:

```
randomForest(formula = term_deposit_factor ~ ., data = train_data,      ntree = 500, importance = TRUE)
```

Type of random forest: classification

Number of trees: 500

No. of variables tried at each split: 4

OOB estimate of error rate: 10.29%

Confusion matrix:

	0	1	class.error
0	25029	566	0.02211369
1	2402	834	0.74227441

```
#confusion matrix of model
```

```
table(Predicted = forrest_pred_class, Actual = test_data$term_deposit_factor)
```

```
      Actual  
Predicted      0      1  
      0 10699  1013  
      1   254   391
```

```
mean(forrest_pred_class == test_data$term_deposit_factor)
```

```
[1] 0.897467
```

```
confusionMatrix(forrest_pred_class, test_data$term_deposit_factor)
```

Confusion Matrix and Statistics

```
      Reference  
Prediction      0      1  
      0 10699  1013  
      1   254   391
```

Accuracy : 0.8975  
95% CI : (0.892, 0.9028)

No Information Rate : 0.8864  
P-Value [Acc > NIR] : 4.377e-05

Kappa : 0.334


Mcnemar's Test P-Value : < 2.2e-16

Sensitivity : 0.9768  
Specificity : 0.2785  
Pos Pred Value : 0.9135  
Neg Pred Value : 0.6062  
Prevalence : 0.8864  
Detection Rate : 0.8658  
Detection Prevalence : 0.9478  
Balanced Accuracy : 0.6277

'Positive' Class : 0

```
varImpPlot(rf_model)
```

rf\_model



## Adaboost

```
#adaboost
ada_partdata <- partdata |> dplyr::select(-y, -term_deposit, -default)
ada_partdata[] <- lapply(ada_partdata, function(x) {
  if (is.character(x)) factor(x) else x
})

set.seed(42)

idx <- createDataPartition(ada_partdata$term_deposit_factor, p = 0.7, list = FALSE)
part_train_data <- ada_partdata[idx, ]
part_test_data <- ada_partdata[-idx, ]

ada_model <- boosting(
```

```

term_deposit_factor ~ .,
data = part_train_data,
mfinal = 100,      # number of boosting iterations
boos = TRUE
)

#ada_model

ada_pred <- predict(ada_model, part_test_data)

ada_pred$class <- factor(
  ada_pred$class,
  levels = levels(part_test_data$term_deposit_factor)
)

confusionMatrix(ada_pred$class, part_test_data$term_deposit_factor)

```

#### Confusion Matrix and Statistics

		Reference	
Prediction	0	1	
0	1053	58	
1	47	77	

Accuracy : 0.915  
 95% CI : (0.898, 0.9299)  
 No Information Rate : 0.8907  
 P-Value [Acc > NIR] : 0.002777  
  
 Kappa : 0.5472  
  
 Mcnemar's Test P-Value : 0.329114  
  
 Sensitivity : 0.9573  
 Specificity : 0.5704  
 Pos Pred Value : 0.9478  
 Neg Pred Value : 0.6210  
 Prevalence : 0.8907  
 Detection Rate : 0.8526  
 Detection Prevalence : 0.8996  
 Balanced Accuracy : 0.7638

```
'Positive' Class : 0
```

```
ada_model$importance
```

	age	campaign	cons.conf.idx	cons.price.idx	contact
10.8225638	4.0182370	1.4636420	1.4424047	0.8964083	
day_of_week	duration	education	emp.var.rate	euribor3m	
5.7047086	26.4282450	7.1719461	1.4904542	10.6128940	
housing	job	loan	marital	month	
2.1332916	10.5441830	0.7079387	2.0421481	9.8731241	
nr.employed	pdays	poutcome	previous		
1.3662911	1.6024729	1.0359175	0.6431296		

## Remove features for all models

remove duration, default, loan, contact, education, employment variation rate. make the number of yes equal to the number of nos for term deposit subscription. Randomly selected term deposit row containing yes to equal the number of term deposit rows containing nos.

```
set.seed(42)
mod_data <- modedata |> dplyr::select( -default, -loan, -contact, -education, -emp.var.rate
#removing highly correlated features

term_deposit_no <- mod_data %>% filter(term_deposit_factor == 0)
term_deposit_yes <- mod_data %>% filter(term_deposit_factor == 1)
#separating the positives and negatives
num_yes <- nrow(term_deposit_yes)

no_sample <- term_deposit_no %>% sample_n(num_yes) #sample the same number of negatives as pos

equal_data <- bind_rows(term_deposit_yes, no_sample) #a dataset with equal number of positive and negative rows

#create the training dataset and test data set based off the equal yes and no data above
eqtrain_index <- sample(seq_len(nrow(equal_data)), size = 0.7 * nrow(equal_data))
eqtrain_data <- equal_data[eqtrain_index, ]
eqtest_data <- equal_data[-eqtrain_index, ]
```

rebuild models

## Rebuild the models

### Decision Trees

```
set.seed(42)
#decision tree model

tree_model <- rpart(
  term_deposit_factor ~ .,
  data = eqtrain_data,
  method = "class",
  control = rpart.control(
    cp = 0.01,
    minsplit = 30,
    maxdepth = 20
  )
)


tree_model
```

n= 6496

```
node), split, n, loss, yval, (yprob)
 * denotes terminal node

1) root 6496 3246 0 (0.5003079 0.4996921)
  2) nr.employed>=5087.65 4699 1678 0 (0.6429027 0.3570973)
     4) month=aug,jul,jun,may,nov 4088 1255 0 (0.6930039 0.3069961) *
     5) month=apr,mar,oct 611 188 1 (0.3076923 0.6923077) *
  3) nr.employed< 5087.65 1797 229 1 (0.1274346 0.8725654) *
```

```
rpart.plot(tree_model)
```



```

pred_class <- predict(tree_model, eqtest_data, type = "class")
#pred_prob <- predict(tree_model, eqtest_data, type = "prob")

table(Predicted = pred_class, Actual = eqtest_data$term_deposit_factor)
  
```

	Actual	
Predicted	0	1
0	1215	566
1	175	828

```
mean(pred_class == eqtest_data$term_deposit_factor)
```

[1] 0.7338362

```
#confusionMatrix(pred_class$class, actual, positive = "1")
confusionMatrix(pred_class, eqtest_data$term_deposit_factor)
```

Confusion Matrix and Statistics

```

    Reference
Prediction   0      1
          0 1215  566
          1 175   828

    Accuracy : 0.7338
    95% CI  : (0.717, 0.7502)
No Information Rate : 0.5007
P-Value [Acc > NIR] : < 2.2e-16

    Kappa : 0.4679

McNemar's Test P-Value : < 2.2e-16

    Sensitivity : 0.8741
    Specificity  : 0.5940
    Pos Pred Value : 0.6822
    Neg Pred Value : 0.8255
    Prevalence    : 0.4993
    Detection Rate : 0.4364
    Detection Prevalence : 0.6397
    Balanced Accuracy : 0.7340

    'Positive' Class : 0

```

0.7338 correct however the data is more balanced.

sensitivity: 0.8741 actual positives , specificity: 0.5940 actual negatives

kappa 0.4679

experiment 1:

cp = 0.001 overfit then prune

```

set.seed(42)
tree_model <- rpart(
  term_deposit_factor ~ .,
  data = eqtrain_data,
  method = "class",
  control = rpart.control(
    cp = 0.001,

```

```

        minsplit = 30,
        maxdepth = 20
    )
)

tree_model

```

n= 6496

```

node), split, n, loss, yval, (yprob)
 * denotes terminal node

```

```

1) root 6496 3246 0 (0.5003079 0.4996921)
2) nr.employed>=5087.65 4699 1678 0 (0.6429027 0.3570973)
   4) month=aug,jul,jun,may,nov 4088 1255 0 (0.6930039 0.3069961)
      8) cons.conf.idx>=-42.35 2430 602 0 (0.7522634 0.2477366) *
      9) cons.conf.idx< -42.35 1658 653 0 (0.6061520 0.3938480)
         18) pdays>=505.5 1601 612 0 (0.6177389 0.3822611)
            36) euribor3m>=1.2755 1307 468 0 (0.6419281 0.3580719)
               72) day_of_week=thu 247 68 0 (0.7246964 0.2753036) *
               73) day_of_week=fri,mon,tue,wed 1060 400 0 (0.6226415 0.3773585)
                  146) previous>=0.5 134 37 0 (0.7238806 0.2761194) *
                  147) previous< 0.5 926 363 0 (0.6079914 0.3920086)
                     294) month=jul 632 232 0 (0.6329114 0.3670886)
                        588) job=blue-collar,housemaid,student,technician,unemployed,unknown 281
                        589) job=admin.,entrepreneur,management,retired,self-employed,services 35
                           1178) euribor3m< 4.9625 283 107 0 (0.6219081 0.3780919) *
                           1179) euribor3m>=4.9625 68 29 1 (0.4264706 0.5735294)
                              2358) job=entrepreneur,retired,services 28 11 0 (0.6071429 0.3928571)
                              2359) job=admin.,management,self-employed 40 12 1 (0.3000000 0.7000000
                           295) month=may 294 131 0 (0.5544218 0.4455782)
                           590) euribor3m< 1.349 263 106 0 (0.5969582 0.4030418)
                              1180) job=admin.,blue-collar,housemaid,services,student 204 73 0 (0.64
                                 1181) job=entrepreneur,management,retired,self-employed,technician,unemp
                                 591) euribor3m>=1.349 31 6 1 (0.1935484 0.8064516) *
                           37) euribor3m< 1.2755 294 144 0 (0.5102041 0.4897959)
                             74) campaign>=5.5 19 2 0 (0.8947368 0.1052632) *
                             75) campaign< 5.5 275 133 1 (0.4836364 0.5163636)
                                150) day_of_week=thu 74 29 0 (0.6081081 0.3918919) *
                                151) day_of_week=fri,mon,tue,wed 201 88 1 (0.4378109 0.5621891)
                                   302) euribor3m< 1.2545 136 65 0 (0.5220588 0.4779412)
                                      604) job=housemaid,management,retired,services,student 31 10 0 (0.67741

```


```

605) job=admin.,blue-collar,entrepreneur,self-employed,technician,unemployed
1210) campaign>=1.5 65    31 0 (0.5230769 0.4769231)
      2420) day_of_week=mon 23     7 0 (0.6956522 0.3043478) *
      2421) day_of_week=fri 42    18 1 (0.4285714 0.5714286) *
      1211) campaign< 1.5 40    16 1 (0.4000000 0.6000000) *
      303) euribor3m>=1.2545 65    17 1 (0.2615385 0.7384615) *
19) pdays< 505.5 57    16 1 (0.2807018 0.7192982) *
5) month=apr,mar,oct 611   188 1 (0.3076923 0.6923077)
10) job=blue-collar,entrepreneur,management,services,technician,unknown 311   139 1 (0
   20) euribor3m>=1.3685 267   133 1 (0.4981273 0.5018727)
   40) euribor3m< 1.4905 200   82 0 (0.5900000 0.4100000)
   80) poutcome=failure 55    15 0 (0.7272727 0.2727273) *
   81) poutcome=nonexistent,success 145   67 0 (0.5379310 0.4620690)
   162) day_of_week=fri,mon 79    30 0 (0.6202532 0.3797468)
   324) marital=divorced,married 58    17 0 (0.7068966 0.2931034) *
   325) marital=single 21     8 1 (0.3809524 0.6190476) *
   163) day_of_week=thu,tue,wed 66    29 1 (0.4393939 0.5606061)
   326) job=entrepreneur,management,services,technician,unknown 44    20 0 (0.5
   327) job=blue-collar 22     5 1 (0.2272727 0.7727273) *
41) euribor3m>=1.4905 67    15 1 (0.2238806 0.7761194) *
21) euribor3m< 1.3685 44     6 1 (0.1363636 0.8636364) *
11) job=admin.,housemaid,retired,self-employed,student,unemployed 300    49 1 (0.16333
3) nr.employed< 5087.65 1797   229 1 (0.1274346 0.8725654) *

```

```
rpart.plot(tree_model)
```

Warning: labs do not fit even at cex 0.15, there may be some overplotting



```
pred_class <- predict(tree_model, eqtest_data, type = "class")
#pred_prob <- predict(tree_model, eqtest_data, type = "prob")
```

```
table(Predicted = pred_class, Actual = eqtest_data$term_deposit_factor)
```

		Actual	
Predicted		0	1
0	1204	518	
1	186	876	

```
mean(pred_class == eqtest_data$term_deposit_factor)
```

```
[1] 0.7471264
```

```
#confusionMatrix(pred_class$class, actual, positive = "1")
confusionMatrix(pred_class, eqtest_data$term_deposit_factor)
```

### Confusion Matrix and Statistics

```

    Reference
Prediction   0      1
          0 1204  518
          1 186   876

    Accuracy : 0.7471
    95% CI  : (0.7305, 0.7632)
No Information Rate : 0.5007
P-Value [Acc > NIR] : < 2.2e-16

    Kappa : 0.4944

McNemar's Test P-Value : < 2.2e-16

    Sensitivity : 0.8662
    Specificity  : 0.6284
    Pos Pred Value : 0.6992
    Neg Pred Value : 0.8249
    Prevalence    : 0.4993
    Detection Rate : 0.4325
    Detection Prevalence : 0.6185
    Balanced Accuracy : 0.7473

    'Positive' Class : 0

```

performs slightly better 0.7471 accuracy

sensitivity: 0.8662

specificity 0.6284

kappa 0.4944

remove additonal features from data: month, day

```

set.seed(42)
feqtrain_data <- eqtrain_data |> dplyr::select(-day_of_week,-month)
feqtest_data <- eqtest_data |> dplyr::select(-day_of_week,-month)

tree_model <- rpart(
  term_deposit_factor ~ .,
  data = feqtrain_data,

```

```

method = "class",
control = rpart.control(
  cp = 0.001,
  minsplit = 30,
  maxdepth = 20
)
)

tree_model

```

n= 6496

```

node), split, n, loss, yval, (yprob)
 * denotes terminal node

1) root 6496 3246 0 (0.50030788 0.49969212)
  2) nr.employed>=5087.65 4699 1678 0 (0.64290275 0.35709725)
    4) cons.conf.idx>=-46.65 4125 1290 0 (0.68727273 0.31272727)
      8) cons.price.idx>=93.956 1400 297 0 (0.78785714 0.21214286) *
      9) cons.price.idx< 93.956 2725 993 0 (0.63559633 0.36440367)
     18) cons.price.idx>=93.0465 1885 622 0 (0.67002653 0.32997347)
       36) euribor3m< 4.207 380 94 0 (0.75263158 0.24736842) *
       37) euribor3m>=4.207 1505 528 0 (0.64916944 0.35083056)
         74) nr.employed>=5211.95 1446 474 0 (0.67219917 0.32780083) *
         75) nr.employed< 5211.95 59 5 1 (0.08474576 0.91525424) *
     19) cons.price.idx< 93.0465 840 371 0 (0.55833333 0.44166667)
       38) euribor3m< 1.349 787 331 0 (0.57941550 0.42058450)
         76) euribor3m>=1.2755 469 170 0 (0.63752665 0.36247335) *
         77) euribor3m< 1.2755 318 157 1 (0.49371069 0.50628931)
           154) campaign>=5.5 19 2 0 (0.89473684 0.10526316) *
           155) campaign< 5.5 299 140 1 (0.46822742 0.53177258)
             310) campaign>=1.5 161 77 0 (0.52173913 0.47826087)
               620) age>=34.5 65 25 0 (0.61538462 0.38461538) *
               621) age< 34.5 96 44 1 (0.45833333 0.54166667)
                 1242) job=admin.,entrepreneur,management,self-employed,technician,unempl
                   2484) age< 28.5 18 4 0 (0.77777778 0.22222222) *
                   2485) age>=28.5 35 15 1 (0.42857143 0.57142857) *
                   1243) job=blue-collar,services,student 43 15 1 (0.34883721 0.65116279)
                     2486) previous>=0.5 16 6 0 (0.62500000 0.37500000) *
                     2487) previous< 0.5 27 5 1 (0.18518519 0.81481481) *
             311) campaign< 1.5 138 56 1 (0.40579710 0.59420290) *
     39) euribor3m>=1.349 53 13 1 (0.24528302 0.75471698) *


```

```

5) cons.conf.idx< -46.65 574 186 1 (0.32404181 0.67595819)
10) job=blue-collar,entrepreneur,management,services,technician,unknown 290 137 1 (0
20) euribor3m>=1.396 243 113 0 (0.53497942 0.46502058)
40) euribor3m< 1.481 190 76 0 (0.60000000 0.40000000)
80) poutcome=failure 54 15 0 (0.72222222 0.27777778) *
81) poutcome=nonexistent,success 136 61 0 (0.55147059 0.44852941)
162) housing=yes 79 30 0 (0.62025316 0.37974684)
324) age>=34.5 48 14 0 (0.70833333 0.29166667) *
325) age< 34.5 31 15 1 (0.48387097 0.51612903)
650) job=services,technician 16 5 0 (0.68750000 0.31250000) *
651) job=blue-collar,entrepreneur,management,unknown 15 4 1 (0.26666666
163) housing=no,unknown 57 26 1 (0.45614035 0.54385965)
326) job=management,unknown 11 3 0 (0.72727273 0.27272727) *
327) job=blue-collar,entrepreneur,services,technician 46 18 1 (0.39130435
41) euribor3m>=1.481 53 16 1 (0.30188679 0.69811321) *
21) euribor3m< 1.396 47 7 1 (0.14893617 0.85106383) *
11) job=admin.,housemaid,retired,self-employed,student,unemployed 284 49 1 (0.1725
3) nr.employed< 5087.65 1797 229 1 (0.12743461 0.87256539) *

```

```
rpart.plot(tree_model)
```



```

pred_class <- predict(tree_model, feqtest_data, type = "class")
#pred_prob <- predict(tree_model, eqtest_data, type = "prob")

table(Predicted = pred_class, Actual = feqtest_data$term_deposit_factor)

```

		Actual
Predicted	0	1
0	1212	523
1	178	871

```
mean(pred_class == feqtest_data$term_deposit_factor)
```

[1] 0.748204

```
#confusionMatrix(pred_class$class, actual, positive = "1")
confusionMatrix(pred_class, feqtest_data$term_deposit_factor)
```

#### Confusion Matrix and Statistics

		Reference
Prediction	0	1
0	1212	523
1	178	871

Accuracy : 0.7482  
95% CI : (0.7316, 0.7642)  
No Information Rate : 0.5007  
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.4966

McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.8719  
Specificity : 0.6248  
Pos Pred Value : 0.6986  
Neg Pred Value : 0.8303  
Prevalence : 0.4993

```
Detection Rate : 0.4353  
Detection Prevalence : 0.6232  
Balanced Accuracy : 0.7484
```

```
'Positive' Class : 0
```

slightly better performance  
accuracy 0.7482  
sensitivity : 0.8719  
specificity: 0.6248 slightly worse  
kappa .4966  
pruning the tree 5

```
set.seed(42)  
tree_model <- rpart(  
  term_deposit_factor ~ .,  
  data = feqtrain_data,  
  method = "class",  
  control = rpart.control(  
    cp = 0.001,  
    minsplit = 30,  
    maxdepth = 5  
  )  
)  
  
tree_model
```

n= 6496


```
node), split, n, loss, yval, (yprob)  
  * denotes terminal node  
  
1) root 6496 3246 0 (0.5003079 0.4996921)  
2) nr.employed>=5087.65 4699 1678 0 (0.6429027 0.3570973)  
4) cons.conf.idx>=-46.65 4125 1290 0 (0.6872727 0.3127273)  
8) cons.price.idx>=93.956 1400 297 0 (0.7878571 0.2121429) *  
9) cons.price.idx< 93.956 2725 993 0 (0.6355963 0.3644037)  
18) cons.price.idx>=93.0465 1885 622 0 (0.6700265 0.3299735) *
```

```

19) cons.price.idx< 93.0465 840 371 0 (0.5583333 0.4416667)
   38) euribor3m< 1.349 787 331 0 (0.5794155 0.4205845) *
   39) euribor3m>=1.349 53 13 1 (0.2452830 0.7547170) *
5) cons.conf.idx< -46.65 574 186 1 (0.3240418 0.6759582)
10) job=blue-collar,entrepreneur,management,services,technician,unknown 290 137 1 (0.4650206)
   20) euribor3m>=1.396 243 113 0 (0.5349794 0.4650206)
      40) euribor3m< 1.481 190 76 0 (0.6000000 0.4000000) *
      41) euribor3m>=1.481 53 16 1 (0.3018868 0.6981132) *
21) euribor3m< 1.396 47 7 1 (0.1489362 0.8510638) *
11) job=admin.,housemaid,retired,self-employed,student,unemployed 284 49 1 (0.1725352)
3) nr.employed< 5087.65 1797 229 1 (0.1274346 0.8725654) *

```

```
rpart.plot(tree_model)
```



```

pred_class <- predict(tree_model, feqtest_data, type = "class")
#pred_prob <- predict(tree_model, eqtest_data, type = "prob")

```

```
table(Predicted = pred_class, Actual = feqtest_data$term_deposit_factor)
```

Actual

```
Predicted      0      1
0 1247  588
1 143   806
```

```
mean(pred_class == feqtest_data$term_deposit_factor)
```

```
[1] 0.7374282
```

```
#confusionMatrix(pred_class$class, actual, positive = "1")
confusionMatrix(pred_class, feqtest_data$term_deposit_factor)
```

#### Confusion Matrix and Statistics

```
Reference
Prediction      0      1
0 1247  588
1 143   806
```

```
Accuracy : 0.7374
95% CI  : (0.7207, 0.7537)
```

```
No Information Rate : 0.5007
P-Value [Acc > NIR] : < 2.2e-16
```

```
Kappa : 0.4751
```

```
McNemar's Test P-Value : < 2.2e-16
```

```
Sensitivity : 0.8971
Specificity  : 0.5782
Pos Pred Value : 0.6796
Neg Pred Value : 0.8493
Prevalence    : 0.4993
Detection Rate : 0.4479
Detection Prevalence : 0.6591
Balanced Accuracy : 0.7377
```

```
'Positive' Class : 0
```

sensitivity improves slightly and everything else is worse.

retraining the Random Forest model with the equally split data

```

set.seed(42)
#random forest model

rf_model <- randomForest(
  term_deposit_factor ~ .,
  data = eqtrain_data,
  #method = "class"
  ntree = 500,
  importance = TRUE
)

forrest_pred_class <- predict(rf_model, eqtest_data, type = "class")
forrest_pred_prob <- predict(rf_model, eqtest_data, type = "prob")

rf_model

```

Call:

```
randomForest(formula = term_deposit_factor ~ ., data = eqtrain_data,      ntree = 500, impo
```

Type of random forest: classification

Number of trees: 500

No. of variables tried at each split: 3

OOB estimate of error rate: 25.03%

Confusion matrix:

	0	1	class.error
0	2830	420	0.1292308
1	1206	2040	0.3715342

```
table(Predicted = forrest_pred_class, Actual = eqtest_data$term_deposit_factor)
```

		Actual
Predicted	0	1
0	1223	544
1	167	850

```
mean(forrest_pred_class == eqtest_data$term_deposit_factor)
```

[1] 0.7446121

```
confusionMatrix(forrest_pred_class, eqtest_data$term_deposit_factor)
```

Confusion Matrix and Statistics

		Reference
Prediction	0	1
0	1223	544
1	167	850

Accuracy : 0.7446  
95% CI : (0.728, 0.7607)  
No Information Rate : 0.5007  
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.4894


McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.8799  
Specificity : 0.6098  
Pos Pred Value : 0.6921  
Neg Pred Value : 0.8358  
Prevalence : 0.4993  
Detection Rate : 0.4393  
Detection Prevalence : 0.6347  
Balanced Accuracy : 0.7448

'Positive' Class : 0

```
varImpPlot(rf_model)
```

## rf\_model



surprisingly model performance is roughly equal to just one decision tree

kappa 0.4894, sensitivity 0.8791, specificity 0.6105 accuracy 0.7446

increase the number of trees from 500 to 1000 (2x)

```
set.seed(42)
rf_model <- randomForest(
  term_deposit_factor ~ .,
  data = eqtrain_data,
  #method = "class"
  ntree = 1000,
  importance = TRUE
)

forrest_pred_class <- predict(rf_model, eqtest_data, type = "class")
forrest_pred_prob <- predict(rf_model, eqtest_data, type = "prob")

rf_model
```

Call:

```
randomForest(formula = term_deposit_factor ~ ., data = eqtrain_data, ntree = 1000, impo
```

```
Number of trees: 1000  
No. of variables tried at each split: 3
```

```
OOB estimate of error rate: 25.18%  
Confusion matrix:  
 0 1 class.error  
0 2825 425 0.1307692  
1 1211 2035 0.3730746
```

```
table(Predicted = forrest_pred_class, Actual = eqtest_data$term_deposit_factor)
```

		Actual
Predicted	0	1
0	1226	542
1	164	852

```
mean(forrest_pred_class == eqtest_data$term_deposit_factor)
```

```
[1] 0.746408
```

```
confusionMatrix(forrest_pred_class, eqtest_data$term_deposit_factor)
```

Confusion Matrix and Statistics

		Reference
Prediction	0	1
0	1226	542
1	164	852

Accuracy : 0.7464  
95% CI : (0.7298, 0.7625)  
No Information Rate : 0.5007  
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.493

Mcnemar's Test P-Value : < 2.2e-16

Sensitivity : 0.8820  
Specificity : 0.6112

```


Pos Pred Value : 0.6934
Neg Pred Value : 0.8386
    Prevalence : 0.4993
Detection Rate : 0.4404
Detection Prevalence : 0.6351
Balanced Accuracy : 0.7466

```

'Positive' Class : 0

```
varImpPlot(rf_model)
```

**rf\_model**



only a slight improvement:

0.745 accuracy,

kappa 0.4901, sensitivity 0.8799, specificity 0.6105.

lets remove features that have both a low mean decrease accuracy and mean decrease gini

day\_of\_week, marital, housing, previous

```

set.seed(42)
rfeqtrain_data <- eqtrain_data |> dplyr::select(-day_of_week,-marital, -housing, -previous)
rfeqtest_data <- eqtest_data |> dplyr::select(-day_of_week,-marital, -housing, -previous)

rf_model <- randomForest(
  term_deposit_factor ~ .,
  data = rfeqtrain_data,
  #method = "class"
  ntree = 1000,
  importance = TRUE
)

forrest_pred_class <- predict(rf_model, rfeqtest_data, type = "class")
forrest_pred_prob <- predict(rf_model, rfeqtest_data, type = "prob")

rf_model

```

Call:

```

randomForest(formula = term_deposit_factor ~ ., data = rfeqtrain_data,      ntree = 1000, in
              Type of random forest: classification
              Number of trees: 1000
              No. of variables tried at each split: 3

              OOB estimate of  error rate: 24.82%
Confusion matrix:
  0     1 class.error
0 2846  404  0.1243077
1 1208 2038  0.3721503

```

```
table(Predicted = forrest_pred_class, Actual = rfeqtest_data$term_deposit_factor)
```

	Actual	
Predicted	0	1
0	1236	556
1	154	838

```
mean(forrest_pred_class == rfeqtest_data$term_deposit_factor)
```

```
[1] 0.7449713
```

```
confusionMatrix(forrest_pred_class, rfeqtest_data$term_deposit_factor)
```

Confusion Matrix and Statistics

		Reference
Prediction	0	1
0	1236	556
1	154	838

Accuracy : 0.745  
95% CI : (0.7283, 0.7611)  
No Information Rate : 0.5007  
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.4902


McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.8892  
Specificity : 0.6011  
Pos Pred Value : 0.6897  
Neg Pred Value : 0.8448  
Prevalence : 0.4993  
Detection Rate : 0.4440  
Detection Prevalence : 0.6437  
Balanced Accuracy : 0.7452

'Positive' Class : 0

```
varImpPlot(rf_model)
```

## rf\_model



Accuracy : 0.7453, Kappa : 0.4909, Sensitivity : 0.8827, Specificity : 0.6083  
specificity decreased

## tuned tree

```
set.seed(42)
rfeqtrain_data$term_deposit_factor <- factor(
  ifelse(rfeqtrain_data$term_deposit_factor == 1, "yes", "no"),
  levels = c("yes", "no")    # convert back to yes and no
)

rfeqtest_data$term_deposit_factor <- factor(
  ifelse(rfeqtest_data$term_deposit_factor == 1, "yes", "no"),
  levels = c("yes", "no")    # convert back to yes and no
)

tuned_rf <- train(
  term_deposit_factor ~ .,
  data = rfeqtrain_data,
  method = "rf",
  trControl = trainControl(
```

```

    method = "cv",
    number = 5,
    classProbs = TRUE,
    summaryFunction = twoClassSummary
),
metric = "ROC",
tuneLength = 5
)

tuned_rf

```

Random Forest

```

6496 samples
10 predictor
2 classes: 'yes', 'no'

```

```

No pre-processing
Resampling: Cross-Validated (5 fold)
Summary of sample sizes: 5197, 5197, 5196, 5197, 5197
Resampling results across tuning parameters:

```

mtry	ROC	Sens	Spec
2	0.7885651	0.6207633	0.8747692
8	0.7951654	0.6512595	0.8403077
15	0.7866247	0.6885348	0.7652308
22	0.7802309	0.6931563	0.7446154
29	0.7780476	0.6956193	0.7390769

```

ROC was used to select the optimal model using the largest value.
The final value used for the model was mtry = 8.

```

```
varImp(tuned_rf)
```

rf variable importance

only 20 most important variables shown (out of 29)

	Overall
euribor3m	100.000
age	66.876

```

nr.employed          63.408
campaign              32.624
cons.conf.idx        17.360
pdays                 15.097
cons.price.idx       14.416
poutcome.success     8.745
poutcome.nonexistent 8.294
jobblue-collar       7.667
jobtechnician         7.347
jobservices           6.258
monthmay               6.098
jobmanagement          5.600
monthoct                4.873
jobself-employed       3.849
jobentrepreneur         3.818
jobretired                3.075
jobunemployed            2.878
jobstudent                  2.839

```

```

forrest_pred_class <- predict(tuned_rf, rfeqtest_data, type = "raw")

table(Predicted = forrest_pred_class, Actual = rfeqtest_data$term_deposit_factor)

```

		Actual	
Predicted	yes	no	
yes	893	197	
no	501	1193	

```
mean(forrest_pred_class == rfeqtest_data$term_deposit_factor)
```

[1] 0.7492816

```
confusionMatrix(forrest_pred_class, rfeqtest_data$term_deposit_factor)
```

#### Confusion Matrix and Statistics

Reference		
Prediction	yes	no
yes	893	197
no	501	1193

```

Accuracy : 0.7493
95% CI  : (0.7327, 0.7653)
No Information Rate : 0.5007
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.4987

McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.6406
Specificity : 0.8583
Pos Pred Value : 0.8193
Neg Pred Value : 0.7043
Prevalence : 0.5007
Detection Rate : 0.3208
Detection Prevalence : 0.3915
Balanced Accuracy : 0.7494

'Positive' Class : yes

```

```
head(tuned_rf$resample)
```

	ROC	Sens	Spec	Resample
1	0.7882103	0.6456086	0.8246154	Fold1
2	0.7893777	0.6379045	0.8261538	Fold5
3	0.7977160	0.6563945	0.8446154	Fold2
4	0.8112722	0.6630769	0.8661538	Fold3
5	0.7892509	0.6533128	0.8400000	Fold4

```
sd(tuned_rf$resample$ROC)
```

```
[1] 0.009782462
```

```
sd(tuned_rf$resample$Kappa)
```

```
[1] NA
```

```

sd(tuned_rf$resample$accuracy)

[1] NA

Accuracy : 0.7478
Kappa : 0.4958
Sensitivity : 0.6413                               Specificity : 0.8547
specificity improved and sensitivity decreased.

#adaptive boosting model

#ada_partdata <- partdata |> dplyr::select(-y, -term_deposit, -default)
#ada_partdata[] <- lapply(ada_partdata, function(x) {
#  if (is.character(x)) factor(x) else x
#})

##ada_partdata$default <- factor(ada_partdata$default, levels = c("no", "yes", "unknown"))

set.seed(42)

#idx <- createDataPartition(ada_partdata$term_deposit_factor, p = 0.7, list = FALSE)
#part_train_data <- ada_partdata[idx, ]
#part_test_data <- ada_partdata[-idx, ]

ada_model <- boosting(
  term_deposit_factor ~ .,
  data = eqtrain_data,
  mfinal = 100,      # number of boosting iterations
  boos = TRUE
)

#ada_model

ada_pred <- predict(ada_model, eqtest_data)

ada_pred$class <- factor(
  ada_pred$class,
  levels = levels(eqtest_data$term_deposit_factor)
)

```

```
confusionMatrix(ada_pred$class, eqtest_data$term_deposit_factor)
```

#### Confusion Matrix and Statistics

```
Reference
Prediction      0      1
      0 1195  503
      1  195  891

Accuracy : 0.7493
95% CI  : (0.7327, 0.7653)
No Information Rate : 0.5007
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.4987

McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.8597
Specificity  : 0.6392
Pos Pred Value : 0.7038
Neg Pred Value : 0.8204
Prevalence   : 0.4993
Detection Rate : 0.4292
Detection Prevalence : 0.6099
Balanced Accuracy : 0.7494

'Positive' Class : 0
```

```
sort(ada_model$importance, decreasing = TRUE)
```

	nr.employed	pdays	euribor3m	month	poutcome
43.1271088	16.3860165	15.7885071	12.5864874	4.8673345	
job	cons.price.idx	age	campaign	day_of_week	
2.1103654	2.0172030	1.2179434	0.5869548	0.5002386	
previous	cons.conf.idx	housing	marital		
0.3902263	0.1722818	0.1329486	0.1163839		

one of the best models so far

Accuracy : 0.7493

Kappa : 0.4987

Sensitivity : 0.8597

Specificity : 0.6392

increase mfinal

```
set.seed(42)
adaeqtrain_data <- eqtrain_data
adaeqtest_data <- eqtest_data

adaeqtrain_data$term_deposit_factor <- factor(
  ifelse(eqtrain_data$term_deposit_factor == 1, "yes", "no"),
  levels = c("yes", "no")
)
adaeqtest_data$term_deposit_factor <- factor(
  ifelse(eqtest_data$term_deposit_factor == 1, "yes", "no"),
  levels = c("yes", "no")
)

ada_model <- boosting(
  term_deposit_factor ~ .,
  data = adaeqtrain_data,
  mfinal = 200,      # number of boosting iterations
  boos = TRUE
)

#ada_model

ada_pred <- predict(ada_model, adaeqtest_data)

ada_pred$class <- factor(
  ada_pred$class,
  levels = levels(adaeqtest_data$term_deposit_factor)
)

confusionMatrix(ada_pred$class, adaeqtest_data$term_deposit_factor)
```

Confusion Matrix and Statistics

```

    Reference
Prediction yes   no
      yes  880 190
      no   514 1200

Accuracy : 0.7471
95% CI : (0.7305, 0.7632)
No Information Rate : 0.5007
P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.4944

McNemar's Test P-Value : < 2.2e-16

Sensitivity : 0.6313
Specificity : 0.8633
Pos Pred Value : 0.8224
Neg Pred Value : 0.7001
Prevalence : 0.5007
Detection Rate : 0.3161
Detection Prevalence : 0.3843
Balanced Accuracy : 0.7473

'Positive' Class : yes

```

```
sort(ada_model$importance, decreasing = TRUE)
```

	nr.employed	euribor3m	pdays	month	poutcome
41.0024894	20.3760138	14.0619534	10.8658134	6.0795533	
job	age	cons.price.idx	day_of_week	campaign	
2.6244126	1.2643830	1.2474492	0.7700321	0.7401658	
previous	cons.conf.idx	housing	marital		
0.4339900	0.2841913	0.1260173	0.1235353		

```
Accuracy : 0.7443 Kappa : 0.4876 Sensitivity : 0.6329
Specificity : 0.8539
```

lets remove features that have both a low importance  
 day\_of\_week, marital, housing, previous

```

set.seed(42)

ada_model <- boosting(
  term_deposit_factor ~ .,
  data = rfeqtrain_data,
  mfinal = 100,      # number of boosting iterations
  boos = TRUE
)

#ada_model

ada_pred <- predict(ada_model, rfeqtest_data)

ada_pred$class <- factor(
  ada_pred$class,
  levels = levels(rfeqtest_data$term_deposit_factor)
)

confusionMatrix(ada_pred$class, rfeqtest_data$term_deposit_factor)

```

#### Confusion Matrix and Statistics

		Reference	
Prediction	yes	no	
yes	875	192	
no	519	1198	

Accuracy : 0.7446  
 95% CI : (0.728, 0.7607)  
 No Information Rate : 0.5007  
 P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.4894

Mcnemar's Test P-Value : < 2.2e-16

Sensitivity : 0.6277  
 Specificity : 0.8619  
 Pos Pred Value : 0.8201  
 Neg Pred Value : 0.6977  
 Prevalence : 0.5007  
 Detection Rate : 0.3143

```
Detection Prevalence : 0.3833  
Balanced Accuracy : 0.7448
```

```
'Positive' Class : yes
```

```
sort(ada_model$importance, decreasing = TRUE)
```

	nr.employed	euribor3m	pdays	month	poutcome
cons.price.idx	31.8924797	31.8622077	16.6109313	8.4684527	4.8096019
	1.8227384	job	age	campaign	cons.conf.idx
	1.7970666		1.5591451	0.8168511	0.3605257

```
Accuracy : 0.7446, Sensitivity : 0.6277, Specificity : 0.8619
```

```
Kappa : 0.4894
```

Model

Changes to data/experiments

Accuracy

Kappa

specificity

sensitivity

note

Decision Tree

Fulldata minus duration

0.8992

0.2667

0.9898

0.1923

Cp = 0.01

Maxdepth = 30

Decision Tree

Remove highly correlated features, made the number of yes equal to the number of no for the term deposits

0.7338

0.4679

0.8741

0.5940

Decision Tree

Same data set as above smaller complexity parameter

0.7471

0.4944

0.8662

0.6284

Cp=0.001

Decision Tree

Removed additional features day\_of\_week and month

0.7482

0.4966

0.8719

0.6248

Cp = 0.001

Decision Tree

The max depth or number of splits is reduced effectively pruning the tree

0.7374

0.4751

0.8971

0.5782

Cp=0.001 maxdepth = 5

Model

Changes to data

Accuracy

Kappa

Specificity

Sensitivity

note

Random Forest

Fulldata minus duration

0.8975

0.334

0.9768

0.2785

Ntree = 500

Random Forest

Remove highly correlated features, made the number of yes equal to the number of no for the term deposits

0.7446

0.4894

0.8799

0.6098

Random Forest

Increased the number of trees in the random forest

0.7464

0.493

0.8820

0.6112

Ntree = 1000

Random Forest

Removed more features based on mean decrease accuracy and mean decrease gini

0.745

0.4902

0.8892

0.6011

Ntree = 1000

Random Forest

5 fold cross validation maximize Area Under the ROC Curve

0.7493

0.4987

0.8583

0.6406

Model

Changes to data

Accuracy

Kappa

specificity

sensitivity

note

adaboost

Started with smaller provided data set due to computation time

0.915

0.5472

0.9573

0.5704

Mfinal = 100

adaboost

Remove highly correlated features, made the number of yes equal to the number of no for the term deposits

0.7493

0.4987

0.8597

0.6392

adaboost

Increase number of boosting iterations

0.7471

0.4944

0.8633

0.6313

Mfinal = 200

Adaboost

Removed features that have a low importance: day\_of\_week, marital, housing, previous

0.7446

0.4894

0.8619

0.6277

The dataset provided contained customer information and campaign information for term deposit subscriptions. Three models were compared. Initially, variance was low when the data was mostly unmodified due to there being a high bias towards non subscribers. The random forest and ensemble methods reduce variance by averaging many trees, while also lowering bias compared to a single tree. AdaBoost reduces bias by focusing on misclassified cases. That was most prominent when comparing the models on a data set that lacked feature engineering.

For the experiments explored in this study sampling of the data and the features that the models were trained on were varied. Also, hyperparameters of the models were adjusted to see the impact on the results. Initially models were biased towards nonsubscribers, so adjustments were made to the models and data to increase accuracy of predicting subscribers.

Initially models were trained on the full data set with 70% of the dataset to be used for training and 30% was used for testing minus duration which was dependent on the target value. All models showed high accuracy and a low kappa. The low kappa signifies that the models were not making predictions much better than statistically random chance. The specificity was high due to the models being biased towards the term deposit not being subscribed to by customers. The number of customers not subscribed heavily outweighed the number of customers that did and the models reflected that. Roughly 10% of the customers subscribed, so if the models said that every customer did not subscribe then the model would be 90% accurate. So, the models showed that they were indeed 90% accurate when the data is 90% not subscribed. The kappa was highest for adaboost and lowest for the decision tree. It should be noted that the adaboost was not originally trained on the full 70% of the dataset due to the computation time. Sensitivity was also in the order of adaboost > random forests > decision tree. This indicated that the decision tree was the worst at predicting subscribers.

To improve the model's sensitivity towards correctly identifying subscribers a new data frame was constructed. The data frame contained all the subscribers and an equal number of randomly sampled nonsubscribers was constructed. The new data frame also had a few highly correlated features removed, such as: default, loan, contact, education, and employment variation rate.

70% of this new 50/50 subscriber/nonsubscriber dataset was used to train the models. They were then tested on the remaining 30% of the dataset. Now an accuracy above 50% would indicate the model may be better than random guessing. When the models were trained on the newly sampled data set the bias shifted towards subscribers more, but it did make the model better at differentiating between the classes since they were equal. The kappa increased for the decision tree and the random forest, indicating that they were not getting their accuracy purely by random chance. Sensitivity, the proportion of true positives correctly identified, also increased for all the models. The overall accuracy and specificity decreased due to the data consisting of significantly less nonsubscribers than the original dataset. The accuracy, kappa, and specificity decreased for adaboost when trained under the new dataset, this may indicate that adaboost works better with complexity or biased data.

For the next set of experiments the models were changed individually. Focusing on the decision tree, it appeared that it may have been underfitted due to only using a couple of features. The complexity parameter was lowered 10-fold to see if a more fitted tree could have more accurate predictions. Accuracy, kappa, and sensitivity all increased. The sensitivity only decreased slightly. The tree had grown significantly more complex by changing this parameter. In order to reduce complexity a couple of more features were removed: day\_of\_week and month. Which had a minimal effect on the model. To avoid any overfitting and reduce complexity the tree was pruned to have a max depth of 5 sets of splits. This resulted in a model that had a lower sensitivity or true positive outcomes. The most complex model had the best results when it came to differentiating subscribers from nonsubscribers.

For the random forest the number of trees in the model increased 2-fold to see if the added complexity would improve the model performance. The model did have a slight increase

in all metrics, however with significantly increased model building time. The mean decrease gini and mean decrease accuracy charts were compared. Features that were both low in decrease in accuracy and decrease in gini (class purity due to split) were then removed from the model: day\_of\_week, marital, housing, and previous. Model performance mostly remained the same after removal of these features, indicating they were not necessary. Finally the random forest was tuned using a 5 fold cross validation. From this there was a decent increase in finding true positives with not much change in accuracy and kappa. Specificity decreased significantly though.

The first unique adaboost experiment was to increase the number of boosting iterations 2-fold which also slightly decreased the sensitivity. Looking at the ranked importance of the features day\_of\_week, marital, housing, and previous were removed. The removal of these features decreased the performance of the model slightly.

It appears that decision trees, adaboost and random forests are very similar in performance especially if the random forest is tuned. It appears that in order to improve the performance of the models additional feature engineering may make a dramatic difference. Since they were trained on a similar and even the same data set the performance is not as dramatic as I would have thought. The decision tree started off as more biased towards nonsubscribers when compared to random forests and adaboost, but its performance improved by making the class sizes equivalent. There are models with very similar metrics between all three. Additional tuning the models would be needed to truly strike a balance between detecting positives and negatives. I would recommend using the random forest due to it achieving the highest sensitivity after tuning. Adaboost is not far behind and may potentially be better with better hyperparameter selection. It appears that the dataset is complex and it models it the unmodified data the best. A complex decision tree performed remarkably similarly to the other two models. With an updating dataset adaboost would be should be better to handle the changes without much revision.

reference:

Input variables: # bank client data: 1 - **age** (numeric) 2 - **job : type of job** (categorical: “admin.”,“blue-collar”,“entrepreneur”,“housemaid”,“management”,“retired”,“self-employed”,“services”,“student”,“technician”,“unemployed”,“unknown”) 3 - **marital : marital status** (categorical: “divorced”,“married”,“single”,“unknown”; note: “divorced” means divorced or widowed) 4 - **education** (categorical: “basic.4y”,“basic.6y”,“basic.9y”,“high.school”,“illiterate”,“profes 5 - **default: has credit in default?** (categorical: “no”,“yes”,“unknown”) 6 - **housing: has housing loan?** (categorical: “no”,“yes”,“unknown”) 7 - **loan: has personal loan?** (categorical: “no”,“yes”,“unknown”) # related with the last contact of the current campaign: 8 - **contact: contact communication type** (categorical: “cellular”,“telephone”) 9 - **month: last contact month of year** (categorical: “jan”, “feb”, “mar”, ..., “nov”, “dec”) 10 - **day\_of\_week: last contact day of the week** (categorical: “mon”,“tue”,“wed”,“thu”,“fri”) 11 - **duration: last contact duration, in seconds (numeric)**. Important note: this attribute highly affects the output target (e.g., if duration=0 then y=“no”). Yet, the duration is not known before a call is performed. Also, after the end

of the call **y** is obviously known. Thus, this input should only be included for benchmark purposes and should be discarded if the intention is to have a realistic predictive model. # other attributes: 12 - **campaign**: number of contacts performed during this campaign and for this client (numeric, includes last contact) 13 - **pdays**: number of days that passed by after the client was last contacted from a previous campaign (numeric; 999 means client was not previously contacted) 14 - **previous**: number of contacts performed before this campaign and for this client (numeric) 15 - **poutcome**: outcome of the previous marketing campaign (categorical: “failure”,“nonexistent”,“success”) # social and economic context attributes 16 - **emp.var.rate**: employment variation rate - quarterly indicator (numeric) 17 - **cons.price.idx**: consumer price index - monthly indicator (numeric) 18 - **cons.conf.idx**: consumer confidence index - monthly indicator (numeric) 19 - **euribor3m**: euribor 3 month rate - daily indicator (numeric) 20 - **nr.employed**: number of employees - quarterly indicator (numeric)

Output variable (desired target): 21 - **y** - has the client subscribed a term deposit? (binary: “yes”,“no”)